

5 Free Courses to Actually Learn AI in 2026

1. Google Machine Learning Crash Course

Best for: complete beginners

Google's own fast-paced intro to ML. Animated videos, interactive visualisations, hands-on exercises. They refreshed it recently to cover modern AI, including a new module on how Large Language Models actually work (tokens, Transformers, the whole thing).

- **What you learn:** linear & logistic regression, classification, neural networks, embeddings, LLMs, production ML
- **Cost:** Free
- **Time:** ~15–20 hours, self-paced
- [🔗 developers.google.com/machine-learning/crash-course](https://developers.google.com/machine-learning/crash-course)

Start here if you've never touched ML before.

2. Hugging Face LLM Course

Best for: anyone who wants to build with LLMs

This is the one if you care about ChatGPT / Claude / Gemini-type models. Built by Hugging Face (the GitHub of AI). Covers Transformers, fine-tuning, and you build real models. There's even a certification at the end.

- **What you learn:** how Transformers work, fine-tuning models, NLP tasks, deploying
 - **Cost:** Free
 - **Time:** Chapter-by-chapter, do it at your pace
 - [🔗 huggingface.co/learn/llm-course](https://huggingface.co/learn/llm-course)
-

3. fast.ai — Practical Deep Learning for Coders

Best for: people who can code and want results fast

Taught by Jeremy Howard (ex-President of Kaggle). The philosophy: build a working model in lesson 1, understand the theory later. Their alumni now work at Google Brain, OpenAI, Adobe, Amazon and Tesla. Over 6 million video views. The full book is free online too.

- **What you learn:** computer vision, NLP, deploying models, building from scratch
- **Prerequisite:** ~1 year of coding (Python preferred). High-school math is enough.
- **Cost:** Free
- **Time:** 9 lessons, ~90 min each
-  course.fast.ai

Best ROI if you already know basic Python.

4. Andrew Ng — Machine Learning Specialization (DeepLearning.AI)

Best for: building strong fundamentals the "proper" way

Andrew Ng is the most famous ML teacher on the planet. This is the updated version of the course that taught millions. Hosted on Coursera. You can **audit it for free** (you only pay if you want the certificate).

- **What you learn:** supervised & unsupervised learning, neural networks, real-world ML best practices
- **Cost:** Free to audit (certificate is paid)
- **Time:** ~2 months at a few hours/week
-  deeplearning.ai/courses/machine-learning-specialization

The "textbook" foundation. Pairs perfectly with fast.ai's hands-on style.

5. Anthropic — Learn (Claude & AI Fluency)

Best for: learning to actually USE AI, not just theory

From the makers of Claude. This is less "train a neural net" and more "become genuinely good at working with AI." Prompt engineering, building with the API, AI fluency. The skills companies are hiring for right now — most job posts ask for hands-on AI users, not researchers.

- **What you learn:** prompting, AI fluency, building with Claude
- **Cost:** Free
-  anthropic.com/learn

The skill that gets you hired in 2026. Don't skip this one.
